

Dimension Reduction and Clustering on the Steam Games Dataset

Elizaveta Sokolova

Contents

Dataset	1
Libraries	1
Data Loading & Column Diagnosis	2
Data Cleaning & Feature Engineering	4
Feature Selection & Log-Transformation	6
Correlation Matrix	7
PCA Validity Test	9
PCA — Principal Component Analysis	9
Clustering	13
Conclusions	15

Dataset

The dataset used in this project is the Steam Games Dataset by FronkonGames, published under a CC-BY-4.0 license. It contains information on over 110,000 games and was collected automatically using the Steam-Games-Scraper — a tool that queries the Steam Web API to retrieve the full app list and then fetches detailed metadata for each title, filtering out DLCs, soundtracks, and utilities so that only actual games are included. The resulting fields cover a wide range of signals: pricing, platform support, review counts, playtime statistics, peak concurrent users, and estimated ownership ranges.

Libraries

```
# Install (run once if packages are missing):
# install.packages(c("FactoMineR", "factoextra", "Rtsne", "umap", "psych",
#                     "ggplot2", "dplyr", "corrplot", "gridExtra", "RColorBrewer", "knitr"))
library(FactoMineR)      # PCA engine (SVD-based)
library(factoextra)        # Scree, biplot, contrib plots (ggplot2-based)
library(psych)             # KMO & Bartlett tests
library(Rtsne)             # t-SNE
library(umap)              # UMAP (pure-R naive method - no Python needed)
library(ggplot2)
library(dplyr)
```

```

library(corrplot)
library(gridExtra)
library(grid)          # textGrob, gpar (used in grid.arrange titles)
library(RColorBrewer)
library(knitr)

```

Data Loading & Column Diagnosis

```

df <- read.csv("games.csv", stringsAsFactors = FALSE)

cat("Dimensions:", nrow(df), "rows  ×", ncol(df), "columns\n\n")

## Dimensions: 122611 rows  × 39 columns

cat("Column names as read from file:\n")

## Column names as read from file:

print(summary(df))

##      AppID           Name        Release.date   Estimated.owners
##  Length:122611  Length:122611  Length:122611    Min.   :0.000e+00
##  Class :character Class :character Class :character  1st Qu.:0.000e+00
##  Mode  :character  Mode  :character  Mode  :character  Median :0.000e+00
##                                         Mean   :5.459e+01
##                                         3rd Qu.:0.000e+00
##                                         Max.   :1.014e+06
##
##      Peak.CCU       Required.age      Price     DiscountDLC.count
##  Min.   : 0.0000  Min.   : 0.000  Min.   : 0.00  Min.   : 0.0000
##  1st Qu.: 0.0000  1st Qu.: 0.550  1st Qu.: 0.00  1st Qu.: 0.0000
##  Median : 0.0000  Median : 2.240  Median : 0.00  Median : 0.0000
##  Mean   : 0.1676  Mean   : 4.765  Mean   :18.35  Mean   : 0.5459
##  3rd Qu.: 0.0000  3rd Qu.: 5.240  3rd Qu.:40.00  3rd Qu.: 0.0000
##  Max.   :21.0000  Max.   :999.980  Max.   :100.00  Max.   :3703.0000
##
##      About.the.game  Supported.languages Full.audio.languages   Reviews
##  Length:122611    Length:122611    Length:122611    Length:122611
##  Class :character Class :character Class :character  Class :character
##  Mode  :character  Mode  :character  Mode  :character  Mode  :character
##
##      Header.image      Website      Support.url      Support.email
##  Length:122611    Length:122611    Length:122611    Length:122611
##  Class :character Class :character Class :character  Class :character

```

```

##  Mode :character  Mode :character  Mode :character  Mode :character
##
##
##
##
##    Windows          Mac          Linux        Metacritic.score
##  Length:122611    Length:122611    Length:122611    Min.   : 0.000
##  Class :character Class :character Class :character  1st Qu.: 0.000
##  Mode  :character Mode  :character Mode  :character  Median  : 0.000
##                                         Mean   : 2.565
##                                         3rd Qu.: 0.000
##                                         Max.   :97.000
##
##  Metacritic.url      User.score       Positive      Negative
##  Length:122611      Min.   : 0.00000  Min.   :     0  Min.   :     0.0
##  Class :character   1st Qu.: 0.00000  1st Qu.:     0  1st Qu.:     0.0
##  Mode  :character   Median : 0.00000  Median :     5  Median :     1.0
##                                         Mean   : 0.02455  Mean   : 1045  Mean   : 169.2
##                                         3rd Qu.: 0.00000  3rd Qu.:    37  3rd Qu.:    10.0
##                                         Max.   :100.00000  Max.   :7642084  Max.   :1173003.0
##
##  Score.rank      Achievements      Recommendations      Notes
##  Min.   : 98.00  Min.   : 0.00  Min.   : 0.0  Length:122611
##  1st Qu.: 99.00  1st Qu.: 0.00  1st Qu.: 0.0  Class :character
##  Median : 99.00  Median : 2.00  Median : 0.0  Mode  :character
##  Mean   : 99.17  Mean   : 18.09  Mean   : 961.8
##  3rd Qu.:100.00  3rd Qu.: 19.00  3rd Qu.: 0.0
##  Max.   :100.00  Max.   :9821.00  Max.   :4830455.0
##  NA's   :122571
##  Average.playtime.forever Average.playtime.two.weeks Median.playtime.forever
##  Min.   :     0      Min.   : 0.00      Min.   :     0.0
##  1st Qu.:     0      1st Qu.: 0.00      1st Qu.:     0.0
##  Median :     0      Median : 0.00      Median :     0.0
##  Mean   : 208      Mean   : 13.79      Mean   : 173.6
##  3rd Qu.:     0      3rd Qu.: 0.00      3rd Qu.:     0.0
##  Max.   :3429544    Max.   :20088.00     Max.   :3429544.0
##
##  Median.playtime.two.weeks Developers      Publishers
##  Min.   : 0.00      Length:122611    Length:122611
##  1st Qu.: 0.00      Class :character Class :character
##  Median : 0.00      Mode  :character Mode  :character
##  Mean   : 14.72
##  3rd Qu.: 0.00
##  Max.   :20088.00
##
##  Categories      Genres          Tags          Screenshots
##  Length:122611    Length:122611    Length:122611    Length:122611
##  Class :character Class :character Class :character Class :character
##  Mode  :character Mode  :character Mode  :character Mode  :character
##
##  Movies

```

```

## Mode:logical
## NA's:122611
##
##
```

Renaming

```

# Rename the first 8 columns to match what the data truly contains
colnames(df) <- gsub("\\.", "_", colnames(df))
colnames(df)
```

```

## [1] "AppID"                  "Name"
## [3] "Release_date"           "Estimated_owners"
## [5] "Peak_CCU"               "Required_age"
## [7] "Price"                  "DiscountDLC_count"
## [9] "About_the_game"          "Supported_languages"
## [11] "Full_audio_languages"   "Reviews"
## [13] "Header_image"           "Website"
## [15] "Support_url"            "Support_email"
## [17] "Windows"                "Mac"
## [19] "Linux"                  "Metacritic_score"
## [21] "Metacritic_url"         "User_score"
## [23] "Positive"               "Negative"
## [25] "Score_rank"              "Achievements"
## [27] "Recommendations"        "Notes"
## [29] "Average_playtime_forever" "Average_playtime_two_weeks"
## [31] "Median_playtime_forever"  "Median_playtime_two_weeks"
## [33] "Developers"              "Publishers"
## [35] "Categories"              "Genres"
## [37] "Tags"                    "Screenshots"
## [39] "Movies"
```

Data Cleaning & Feature Engineering

Parse string columns

```

# Estimated_owners: "20000 - 50000" + midpoint 35000
parse_owners <- function(x) {
  x <- gsub(", ", "", as.character(x))

  parts <- strsplit(x, "\\s*[-]\\s*")

  sapply(parts, function(p) {
    nums <- suppressWarnings(as.numeric(p))
    # ,
```

```

    if (any(!is.na(nums))) mean(nums, na.rm = TRUE) else NA_real_
  }, USE.NAMES = FALSE)
}

df$Owners_numeric <- parse_owners(df$Estimated_owners)

cat("Owners_numeric - sample values:\n")

## Owners_numeric - sample values:

df$Owners_numeric <- parse_owners(df$Estimated_owners)
head(df$Owners_numeric, 10)

## [1] 0 0 0 1 0 0 8 0 0 1

```

Filter & group genres

```

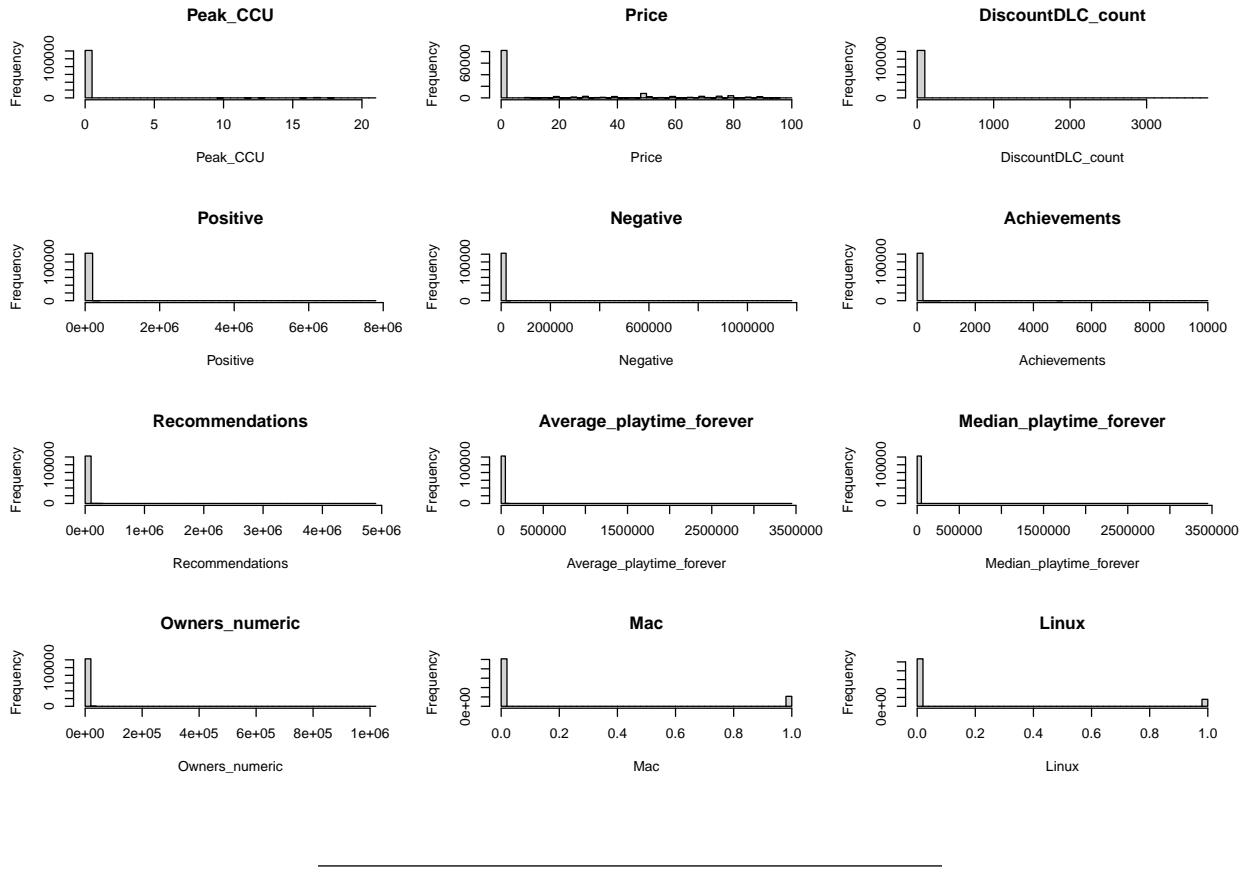
par(mfrow = c(4, 3))
keep_cols <- c("Peak_CCU", "Price", "DiscountDLC_count",
              "Positive", "Negative", "Achievements", "Recommendations",
              "Average_playtime_forever", "Median_playtime_forever",
              "Owners_numeric", "Mac", "Linux")

feat <- df[, keep_cols]

# Binary platform flags
feat$Mac <- as.integer(as.logical(feat$Mac))
feat$Linux <- as.integer(as.logical(feat$Linux))

for(col in keep_cols) {
  hist(feat[[col]], breaks = 50, main = col, xlab = col)
}

```



Feature Selection & Log-Transformation

Why these features?

Feature	Transform	Rationale
Peak CCU	log	Peak concurrent users — extreme right skew
Price	none	Pricing; mild skew
DLC Count	log	Ecosystem depth
Positive Reviews	log	Core engagement signal
Negative Reviews	log	Quality / controversy
Achievements	log	Game-design depth
Recommendations	log	Community endorsement
Avg Playtime	log	Player retention
Med Playtime	log	Retention (outlier-robust)
Est. Owners	log	Market reach
Mac	none	Binary platform flag
Linux	none	Binary platform flag

Dropped: * Required_age — 98.6 % zeros (near-zero variance). * Metacritic score — 94.7 % zeros. * Windows — TRUE for 100 % of games (zero variance).

```

# log on all skewed continuous columns
skewed <- c("Peak_CCU", "Positive", "Negative", "Recommendations",
           "Average_playtime_forever", "Median_playtime_forever",
           "Owners_numeric", "DiscountDLC_count", "Achievements")

feat[, skewed] <- log1p(feat[, skewed])

# Clean column names for nicer plots
colnames(feat) <- c("Peak CCU", "Price", "DLC Count",
                     "Positive", "Negative", "Achievements", "Recommendations",
                     "Avg Playtime", "Med Playtime", "Est. Owners",
                     "Mac", "Linux")

cat("Final feature matrix:", nrow(feat), "games  x", ncol(feat), "features\n\n")

## Final feature matrix: 122611 games  x 12 features

print(summary(feat))

```

```

##      Peak CCU          Price        DLC Count       Positive
##  Min. :0.00000  Min.   : 0.00  Min.   :0.0000  Min.   : 0.000
##  1st Qu.:0.00000 1st Qu.: 0.00  1st Qu.:0.0000  1st Qu.: 0.000
##  Median :0.00000 Median : 0.00  Median :0.0000  Median : 1.792
##  Mean   :0.02935 Mean   :18.35  Mean   :0.1485  Mean   : 2.266
##  3rd Qu.:0.00000 3rd Qu.:40.00  3rd Qu.:0.0000  3rd Qu.: 3.638
##  Max.   :3.09104 Max.   :100.00  Max.   :8.2172  Max.   :15.849
##      Negative        Achievements  Recommendations  Avg Playtime
##  Min.   : 0.0000  Min.   :0.000  Min.   : 0.000  Min.   : 0.000
##  1st Qu.: 0.0000  1st Qu.:0.000  1st Qu.: 0.000  1st Qu.: 0.000
##  Median : 0.6931  Median :1.099  Median : 0.000  Median : 0.000
##  Mean   : 1.4764  Mean   :1.512  Mean   : 1.099  Mean   : 1.121
##  3rd Qu.: 2.3979  3rd Qu.:2.996  3rd Qu.: 0.000  3rd Qu.: 0.000
##  Max.   :13.9751  Max.   :9.192  Max.   :15.390  Max.   :15.048
##      Med Playtime    Est. Owners          Mac          Linux
##  Min.   : 0.000  Min.   :0.00000  Min.   :0.0000  Min.   :0.0000
##  1st Qu.: 0.000  1st Qu.:0.00000  1st Qu.:0.0000  1st Qu.:0.0000
##  Median : 0.000  Median : 0.0000  Median :0.0000  Median :0.0000
##  Mean   : 1.093  Mean   : 0.3333  Mean   :0.1737  Mean   :0.1281
##  3rd Qu.: 0.000  3rd Qu.:0.0000  3rd Qu.:0.0000  3rd Qu.:0.0000
##  Max.   :15.048  Max.   :13.8294  Max.   :1.0000  Max.   :1.0000

```

Correlation Matrix

```

cor_mat <- cor(feat, use = "pairwise.complete.obs")

corrplot(cor_mat,
         method   = "color",
         col      = colorRampPalette(c("#2166ac", "#f7f7f7", "#b2182b"))(200),

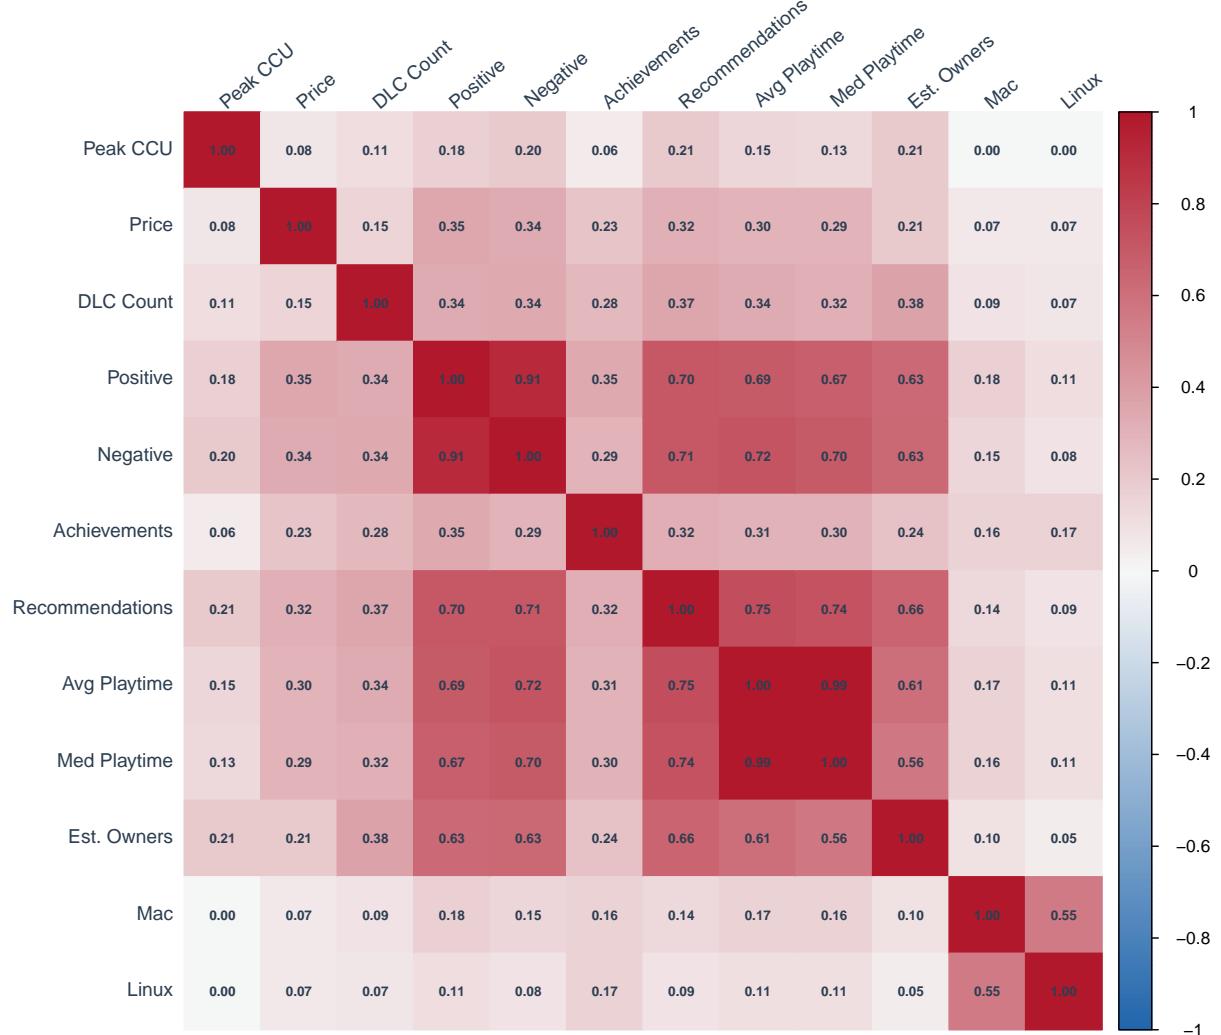
```

```

tl.cex      = 0.82,
tl.col       = "#2c3e50",
tl.srt       = 40,
addCoef.col = "#2c3e50",
number.cex   = 0.60,
number.digits = 2,
cl.pos      = "r",
cl.cex       = 0.72,
title        = "Pairwise Correlations - Steam Game Features\n(log-transformed where indicated"
mar         = c(0, 0, 2, 0))

```

**Pairwise Correlations — Steam Game Features
(log-transformed where indicated)**



Several feature groups are strongly correlated — *Positive / Recommendations / Est. Owners / Avg & Med Playtime / Peak CCU* form one tight block; *Avg Playtime* and *Med Playtime* are near-duplicates. This redundancy is exactly what PCA is designed to compress.

PCA Validity Test

Before running PCA we should verify that the correlation structure in the data is actually suitable for factor extraction. Two standard tests exist for this:

- **Kaiser-Meyer-Olkin (KMO)** — measures *sampling adequacy*. It checks whether the pairwise correlations are large relative to the partial correlations (i.e., whether the variables share enough common variance to justify reducing them). KMO > 0.6 is acceptable; > 0.8 is good; > 0.9 is excellent.

```
# KMO needs the raw data (not a cor matrix) to compute per-variable MSA.  
kmo_res      <- KMO(feat)  
print(kmo_res)
```

```
## Kaiser-Meyer-Olkin factor adequacy  
## Call: KMO(r = feat)  
## Overall MSA =  0.81  
## MSA for each item =  
##          Peak CCU          Price        DLC Count       Positive       Negative  
##          0.91          0.96          0.94          0.82          0.83  
## Achievements Recommendations Avg Playtime Med Playtime Est. Owners  
##          0.89          0.96          0.74          0.73          0.85  
##          Mac           Linux  
##          0.62          0.57
```

PCA — Principal Component Analysis

Run PCA

```
# FactoMineR's PCA() centres and scales every variable to unit variance automatically.  
res.pca <- PCA(feat, graph = FALSE, scale.unit = TRUE)
```

```
# Eigenvalue table - as.data.frame so that $ column access works  
eigen_tbl <- as.data.frame(get_eigenvalue(res.pca))
```

```
cat("Eigenvalues & explained variance (all components):\n\n")
```

```
## Eigenvalues & explained variance (all components):
```

```
print(round(eigen_tbl, 3))
```

```
##          eigenvalue variance.percent cumulative.variance.percent  
## Dim.1      5.186        43.217            43.217  
## Dim.2      1.534        12.782            56.000  
## Dim.3      0.965        8.039            64.039  
## Dim.4      0.943        7.862            71.901  
## Dim.5      0.869        7.239            79.141  
## Dim.6      0.695        5.792            84.932
```

```

## Dim.7      0.568      4.733      89.666
## Dim.8      0.448      3.731      93.397
## Dim.9      0.428      3.566      96.962
## Dim.10     0.273      2.276      99.238
## Dim.11     0.086      0.714      99.952
## Dim.12     0.006      0.048      100.000

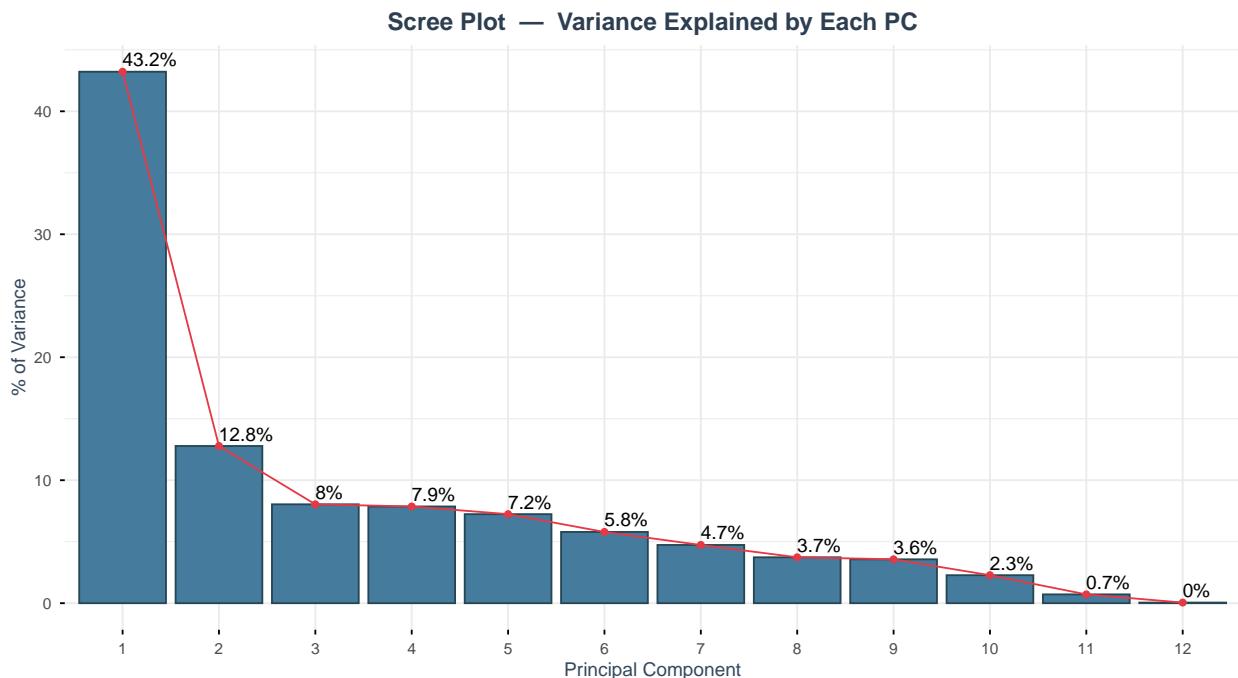
```

Scree Plot

```

fviz_eig(res.pca,
          addlabels = TRUE,
          ncp      = ncol(feat),
          barfill  = "#457b9d",
          barcolor  = "#264653",
          linecolor = "#e63946",
          pointcolor = "#e63946",
          pointsize = 3.8,
          ggtheme   = theme_minimal() +
  labs(title = "Scree Plot - Variance Explained by Each PC",
       x      = "Principal Component",
       y      = "% of Variance") +
  theme(plot.title = element_text(face = "bold", size = 14,
                                  hjust = 0.5, color = "#2c3e50"),
        axis.title = element_text(size = 11, color = "#34495e"))

```



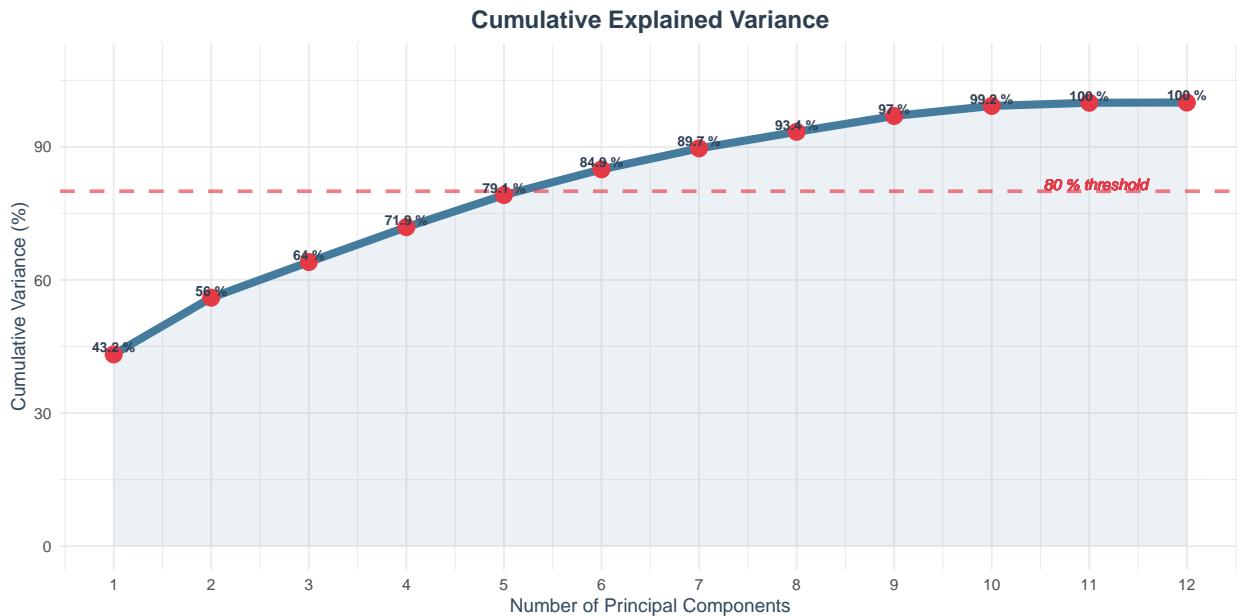
Cumulative Explained Variance

```

cum_var <- eigen_tbl[, 3]    # col 3 = cumulative percentage of variance
n_pc      <- length(cum_var)

ggplot(data.frame(PC = 1:n_pc, CumVar = cum_var)) +
  geom_area(aes(x = PC, y = CumVar), fill = "#457b9d", alpha = 0.10) +
  geom_line(aes(x = PC, y = CumVar), color = "#457b9d", linewidth = 2.2) +
  geom_point(aes(x = PC, y = CumVar), color = "#e63946", size = 4.2) +
  # 80 % threshold
  geom_hline(yintercept = 80, linetype = "dashed",
             color = "#e63946", linewidth = 1, alpha = 0.65) +
  geom_text(x = n_pc - 0.4, y = 81.5,
            label = "80 % threshold", hjust = 1,
            color = "#e63946", size = 3.1, fontface = "italic") +
  # value labels
  geom_text(aes(x = PC, y = CumVar + 1.6),
            label = paste0(round(cum_var, 1), " %"),
            hjust = 0.5, size = 2.8, color = "#2c3e50", fontface = "bold") +
  scale_x_continuous(breaks = 1:n_pc) +
  scale_y_continuous(limits = c(0, 108)) +
  labs(title = "Cumulative Explained Variance",
       x      = "Number of Principal Components",
       y      = "Cumulative Variance (%)") +
  theme_minimal() +
  theme(plot.title = element_text(face = "bold", size = 14,
                                  hjust = 0.5, color = "#2c3e50"),
        axis.title = element_text(size = 11, color = "#34495e"))

```



Variable Contributions to PC1 & PC2

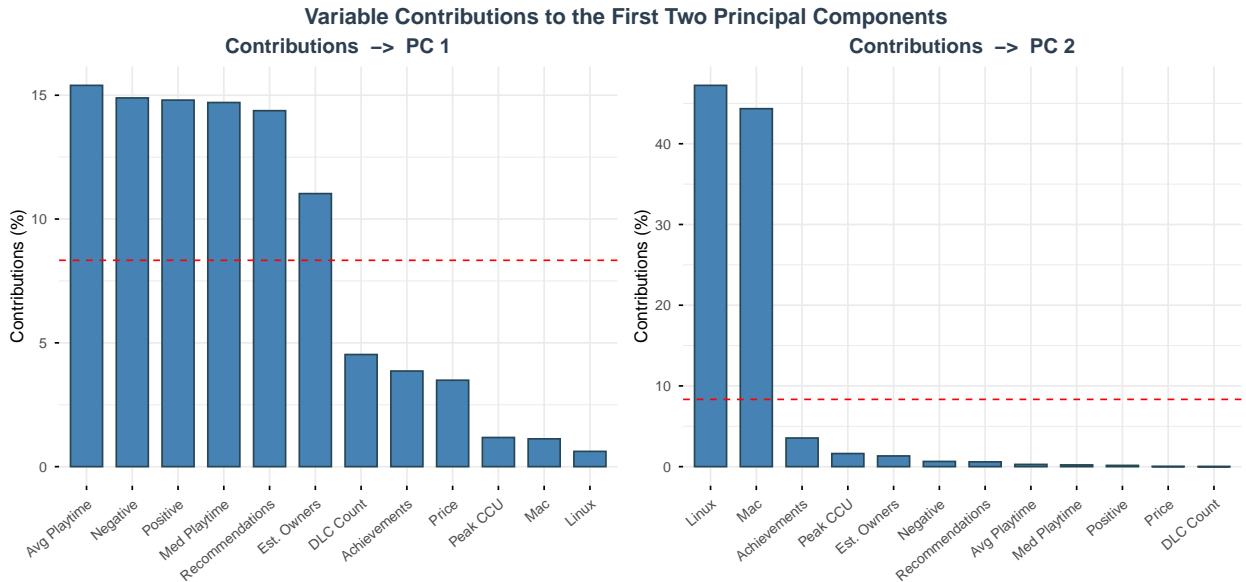
```

p1 <- fviz_contrib(res.pca, choice = "var", axes = 1,
                   fill.color = "#457b9d", color = "#264653") +
  labs(title = "Contributions → PC 1") +
  theme(axis.text.x = element_text(angle = 40, hjust = 1, size = 9),
        plot.title = element_text(face = "bold", size = 13,
                                  hjust = 0.5, color = "#2c3e50"))

p2 <- fviz_contrib(res.pca, choice = "var", axes = 2,
                   fill.color = "#2a9d8f", color = "#264653") +
  labs(title = "Contributions → PC 2") +
  theme(axis.text.x = element_text(angle = 40, hjust = 1, size = 9),
        plot.title = element_text(face = "bold", size = 13,
                                  hjust = 0.5, color = "#2c3e50"))

grid.arrange(p1, p2, ncol = 2,
             top = textGrob("Variable Contributions to the First Two Principal Components",
                            gp = gpar(fontsize = 14, fontface = "bold", col = "#2c3e50")))

```



Variable Plot (Correlation Circle)

```

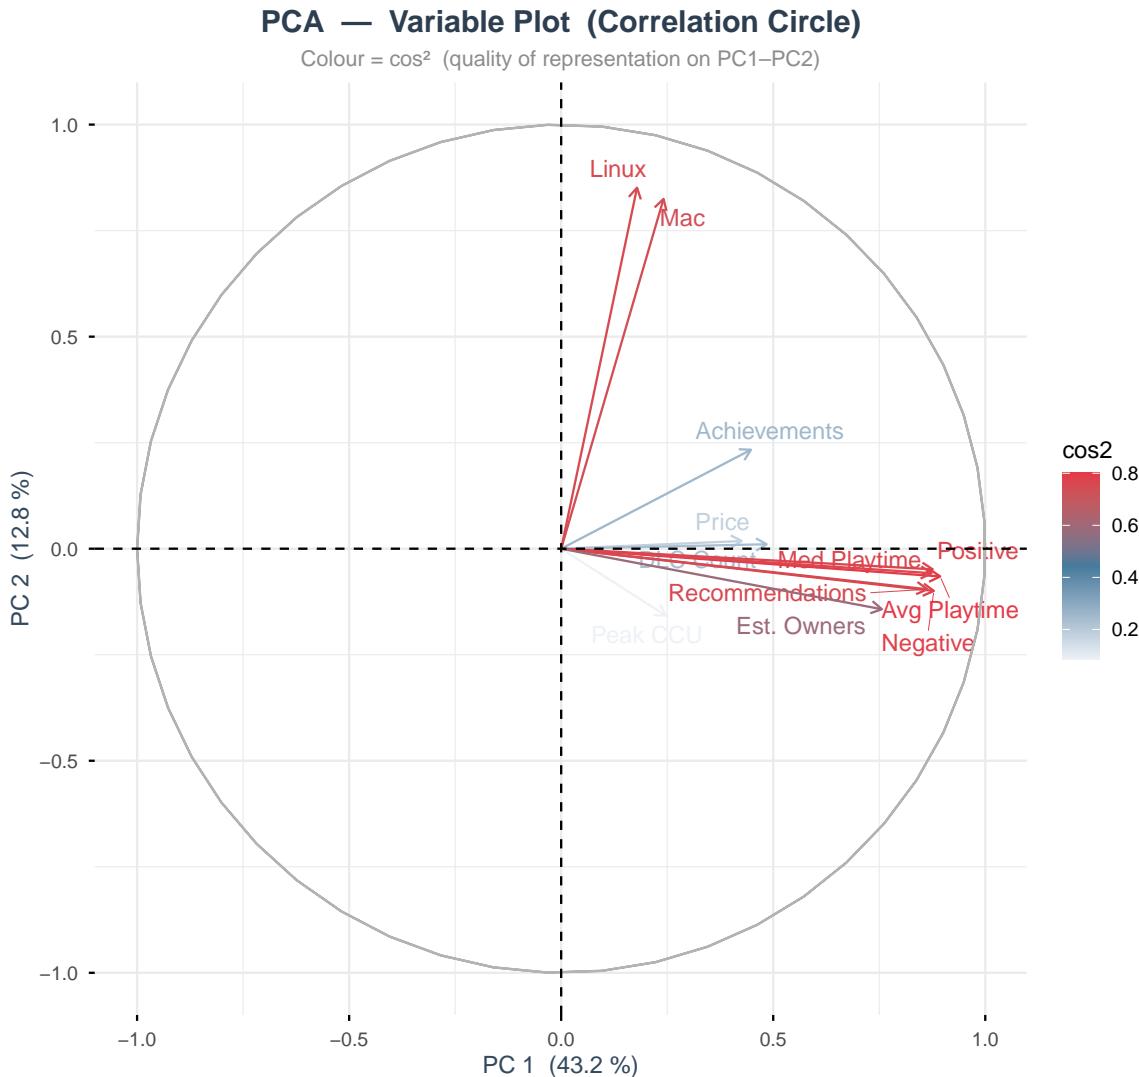
fviz_pca_var(res.pca,
              col.var      = "cos2",
              gradient.cols = list("white" = "#eef2f7",
                                    "blue"   = "#457b9d",
                                    "red"    = "#e63946"),
              repel       = TRUE,
              geom.var    = c("arrow", "text"),
              ggtheme     = theme_minimal(),
              legend.position = "bottom") +
  labs(title = "PCA - Variable Plot (Correlation Circle)",
       subtitle = "Colour = cos² (quality of representation on PC1-PC2)",

```

```

x = paste0("PC 1  (", round(eigen_tbl[1, 2], 1), " %)"),
y = paste0("PC 2  (", round(eigen_tbl[2, 2], 1), " %)") + 
theme(plot.title    = element_text(face = "bold", size = 14,
                                    hjust = 0.5, color = "#2c3e50"),
      plot.subtitle = element_text(size = 10, color = "grey55", hjust = 0.5),
      axis.title    = element_text(size = 11, color = "#34495e"))

```



Clustering

After considering K-means, hierarchical, and DBSCAN clustering, DBSCAN was chosen as the final method due to its ability to handle skewed engagement metrics and uneven density distribution without assuming spherical cluster shapes.

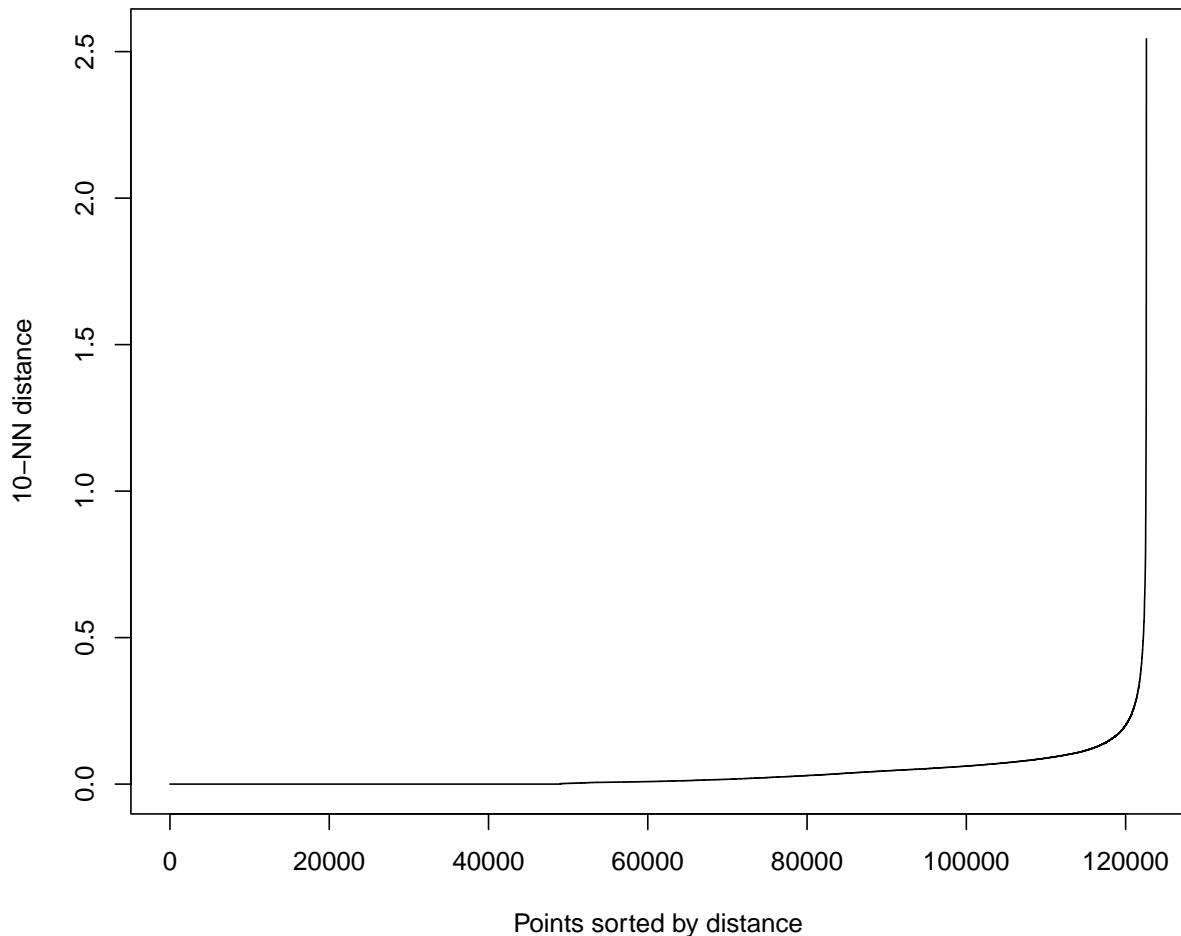
```

library(dbSCAN)
pca_scores <- as.data.frame(res.pca$ind$coord)

# Lets take first 3 components

```

```
pca_sub <- scale(pca_scores[,1:3])  
kNNdistplot(pca_sub, k = 10)
```



```
db <- dbSCAN(pca_sub, eps = 0.28, minPts = 30)  
df$Cluster_DBSCAN <- factor(db$cluster)  
table(df$Cluster_DBSCAN)
```

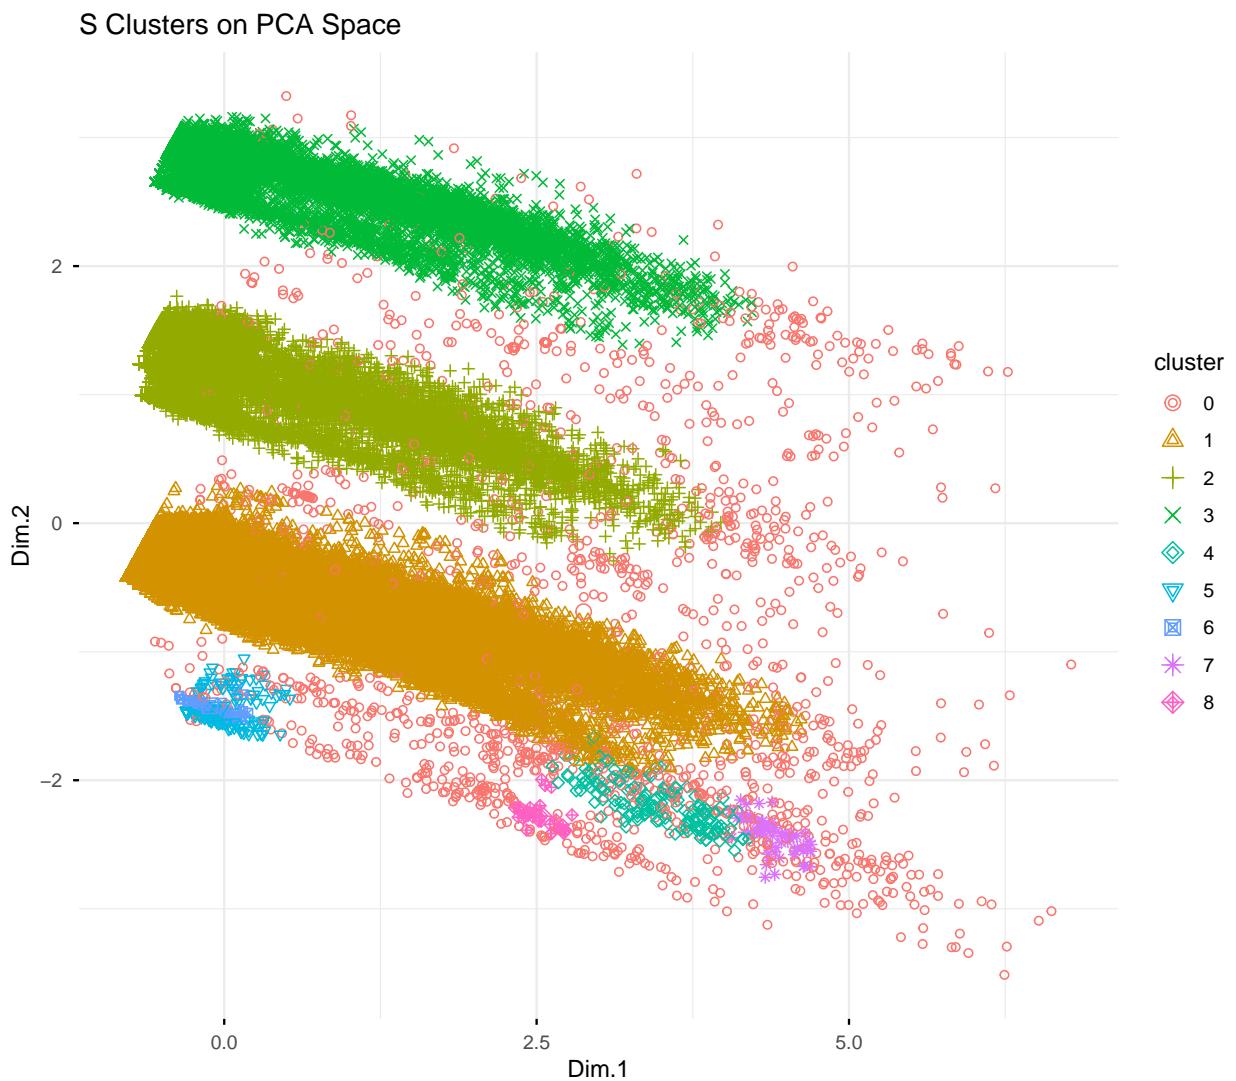
```
##  
##      0      1      2      3      4      5      6      7      8  
## 1614 95513 14022 10994   164    160    47    61    36
```

```
fviz_cluster(list(data = pca_sub[,1:2],  
                  cluster = db$cluster),  
            geom = "point",
```

```

ellipse = FALSE,
ggtheme = theme_minimal(),
main = "S Clusters on PCA Space")

```



Conclusions

PCA

Principal Component Analysis revealed a strong dimensional structure in the dataset.

PC1 is mainly driven by popularity / engagement variables — Peak CCU, Positive Reviews, Recommendations, Estimated Owners, Playtime. This component can be interpreted as a “Market Success & Player Engagement” axis.

PC2 is influenced more by game design / ecosystem features

Clustering

Density-based clustering (DBSCAN) revealed natural groupings of games and isolated outliers without assuming spherical cluster shapes like K-means has due to highly skewed engagement metrics and uneven density distribution in the Steam marketplace. Most games ended up in one large group with similar low engagement.